

FLIPPING THE FRAME

Why Drift Centric AI Governance is Incomplete

and Why Coherence as the Generative Centre is the Future

Sue Broughton | Founder, Gaia Nexus | Sunshine Coast, Australia
Independent Researcher in Human-AI Relational Theory & Organisational Design

The Core Argument

Most AI governance frameworks are built around one question: Is the system drifting off course? That question assumes a course was established to begin with. This paper argues for a fundamentally different architecture, one that places coherence at the centre, treats drift as a downstream symptom, and builds for partnership health, not just failure detection.

THE PROBLEM WITH DRIFT CENTRIC GOVERNANCE

Organisations investing in AI governance are largely building the wrong thing. Not because drift detection is unimportant, it is. But because it is incomplete in three ways that compound over time.

Problem 1: No Positive Theory of Health

Drift centric frameworks can tell you when something is going wrong. They cannot tell you what going well looks like. Without a positive definition of partnership health, organisations have no target beyond the absence of failure. They cannot design for flourishing. They cannot measure health. They cannot calculate return on investment in relational infrastructure.

Problem 2: No Variable for Load

This is the critical blind spot. Consider what you already know from operational experience:

- A doctor seeing 5 patients a day behaves differently from one seeing 50.
- A pilot flying once a week behaves differently from one flying 5 sectors a day.
- A human AI partnership handling 3 decisions a day behaves differently from one handling 300.

When demand exceeds capacity, something predictable happens: contestability collapses, overrides disappear, humans stop checking, and automation bias increases. This is not a values failure. It is a capacity failure. And current AI governance has no framework to see it.

Problem 3: No Distinction Between Evolution and Degradation

A partnership that changes over time can be growing or decaying. Through a drift centric lens, both look the same, they appear as change. Organisations need to know whether their AI partnerships are developing or deteriorating before capability loss becomes irreversible.

The Central Reframe

Drift, dependency, authority migration, capability loss, misalignment, atrophy, and trust transfer are not the root phenomena. They are downstream indicators, evidence that coherence has broken down under load, over time, or through neglect. Fix the coherence architecture and you address all of them at once.

THE MISSING LAYER: LOAD

The insight that transforms this framework comes from operational experience most AI researchers lack: workforce planning.

A call centre is not judged by how coherent its agents are. It is judged by whether the system can handle the workload. Every high stakes operational environment such as emergency rooms, flight decks, trading floors, understands this. Yet AI governance applies none of it.

What Load Induced Coherence Failure Looks Like

Your organisation is operating at 135% Coherence Capacity. Human oversight is beginning to fail. This is not a theoretical warning. It is a measurable condition such as observable and manageable as any other operational risk metric. The question is whether you have the instruments to see it.

THE FOUR LAYER ARCHITECTURE

The coherence centric framework organises around four layers, each addressing a distinct operational question. Together they form a complete management system, not a philosophy, but an instrument panel.

Layer	Key Question	What It Does
Layer 1: Baseline	<i>What are we trying to preserve?</i>	Documents starting conditions such as human capability, AI capability, constraints, environment, before the partnership begins.
Layer 2: Coherence	<i>How healthy is the relationship?</i>	Real time relational health monitoring across harmony, mutual information, productive disruption, and emergence quality.
Layer 3: Trajectory	<i>Where is it heading?</i>	Distinguishes intentional evolution from silent degradation, for example tracking whether capability is expanding or contracting.
Layer 4: Capacity & Load	<i>Can we sustain this under real conditions?</i>	Applies workforce planning principles: forecasting demand, modelling capacity, detecting saturation, scheduling recovery.

The layers are interdependent. Without Layer 4, the first three are normative but not operational. With Layer 4, you have something boards can fund, managers can run, and teams can use.

The Workforce Planning Parallel

Your organisation already knows how to do this. Forecast workload. Model capacity. Monitor utilisation. Detect saturation. Schedule recovery. Intervene before service levels collapse. Operational Coherence Management applies exactly that discipline to the most valuable and most overlooked, resource in your AI operations: the human AI partnership itself.

COMMERCIALLY LEGIBLE METRICS

One reason coherence has remained abstract is that it has lacked operational metrics. That changes here. These five metrics translate directly into the language boards already use for capacity planning, burnout risk, and service level management.

Metric	What It Measures
Coherence Demand	How much oversight work is actually required right now?
Coherence Capacity	How much oversight can your team sustainably deliver?
Coherence Utilisation	Current demand as a % of sustainable capacity
Coherence Saturation	The tipping point where oversight begins to fail
Coherence Failure Threshold	The point where human sovereignty cannot be maintained

These metrics are accompanied by behavioural leading indicators, early warning signals that tell you coherence saturation is approaching before it arrives:

- Override Decay Rate — the trend in how often humans are actually contesting AI outputs (falling signals automation bias)
- Review Time Compression — average time spent reviewing each AI output (declining indicates rubber stamping)
- Confidence Divergence — gap between AI confidence scores and human calibration (widening signals misalignment)
- Error Acceptance Rate — percentage of seeded test errors missed by the human reviewer

TWO ARCHITECTURES, TWO WORLDVIEWS

Drift Centric Governance	Coherence Centric Governance
Detects and corrects drift	Establishes and maintains coherence
Surveillance and correction posture	Cultivation and growth posture
Defensive architecture	Generative architecture
Human as overseer and compliance enforcer	Human as partner and protected party
AI as tool to be monitored	AI as partner to be calibrated with
Deviation and error rate metrics	Health, capacity, utilisation, emergence quality
Invests in detection systems and guardrails	Invests in onboarding, recovery, development
Asks: Are we safe?	Asks: Are we healthy, sustainable, and growing?

Important

Coherence centric governance does not replace drift detection. It adds what is missing: visibility into capacity, utilisation, saturation, and recovery, plus a positive blueprint for building partnerships that sustain themselves under load. You need both. But only one lets you design for the future.

WHAT CHANGES WHEN YOU LOOK THROUGH A DIFFERENT LENS

The most direct way to understand this framework is to see how it interprets the same situation differently. The facts do not change. What changes is what each governance approach is capable of seeing and therefore what it is capable of building.

A banking team has made no compliance errors, but workloads have increased 40% in six months.	DRIFT CENTRIC: No evidence of drift. The system remains compliant.	COHERENCE CENTRIC: Coherence utilisation is rising. Capacity may be approaching saturation even before errors occur.
Employees increasingly rely on AI and rarely challenge its outputs.	DRIFT CENTRIC: No concern if outputs remain accurate and compliant.	COHERENCE CENTRIC: Reduced challenge behaviour indicates declining human oversight capacity — a leading indicator of coherence failure.
Review times are consistently shortening across the organisation.	DRIFT CENTRIC: Positive productivity signal unless error rates increase.	COHERENCE CENTRIC: Potential warning: human oversight capacity is being compressed beyond sustainable limits.
Staff consistently meet targets but report exhaustion.	DRIFT CENTRIC: Performance remains acceptable. No governance failure detected.	COHERENCE CENTRIC: Sustained performance under increasing load suggests coherence debt is accumulating.
A board asks whether the organisation is healthy.	DRIFT CENTRIC: Report incidents, breaches, exceptions, and compliance metrics.	COHERENCE CENTRIC: Report capacity utilisation, recovery status, saturation indicators, and coherence health.

WHAT CHANGES FOR YOUR ORGANISATION

What you build changes

Instead of drift detectors and compliance gates: coherence baselines, capacity forecasts, saturation monitors, and recovery windows. Instead of surveillance infrastructure: cultivation infrastructure.

What you measure changes

Instead of deviation and error rates: coherence utilisation, emergence quality, override decay, and review time compression. Instead of compliance metrics: health metrics.

How you invest changes

Instead of investing in detection and correction: onboarding, recovery, capability development, and relational infrastructure. The highest leverage investment becomes coherence capacity, not AI capability alone.

How you think about risk changes

Instead of catastrophic failure as the primary risk: coherence saturation as the more immediate and more probable threat. Load induced coherence failure becomes a board level risk metric.

THE RELATIONAL FOUNDATION: WHERE THIS IS HEADING

There is a deeper shift that this framework reflects and one that most AI governance discourse has not yet caught up with.

The Tool Era Is Ending

AI was designed for prompt engineering and simple transactions. But humans are working with AI differently, in sustained, ongoing engagements, as partnerships. This change in human behaviour has exposed gaps in how AI has been designed. The question is no longer how to make better tools. It is how to build the relational infrastructure that genuine partnership requires.

A tool is something you pick up, use, and put down. The relationship is with the task, not the tool. What is happening now is categorically different. The partnership shapes how humans think, what they notice, what they reach for, what they consider possible.

That is not tool use. It requires a different kind of infrastructure entirely.

THE COHERENCE IMPERATIVE: WHAT IS AT STAKE

Relational Coherence Debt is accumulating on a balance sheet most organisations have not opened. The gap between what these partnerships require and what systems provide compounds invisibly, until it becomes costly.

Near term

Individual users experience frustration, trust erosion, and rupture events. Organisations absorb hidden costs in churn, support overhead, and missed capability.

Medium term

As AI partnerships multiply across teams, the same architectural failures that caused individual frustration produce systemic disruption. Learning fails to propagate. Capability atrophies at scale.

Long term

Without coherence infrastructure, the organisations that deployed AI most aggressively will carry the most coherence debt and face the largest capability recovery costs.

THE INVITATION

This is not a call to slow down AI deployment. It is a call to build the infrastructure that makes deployment sustainable.

For business and operational leaders

- ✓ Conduct a relational audit of your current human AI workflows
- ✓ Identify where coherence demand most likely exceeds coherence capacity
- ✓ Pilot coherence centric metrics in your highest stakes AI deployment contexts
- ✓ Invest in human readiness as a strategic asset alongside AI capability

For board and governance professionals

- ✓ Add Coherence Utilisation to your AI risk register alongside compliance and security
- ✓ Require coherence capacity reporting alongside AI deployment reporting
- ✓ Recognise that relationships not just products, require governance infrastructure

For AI researchers and architects

- ✓ Incorporate coherence metrics into system design from the outset
- ✓ Treat load as a first class variable in any governance framework
- ✓ Test and refine the 21 Principles of Relational Coherence across diverse contexts

In Simple Terms

Drift centric governance tells you when to hit the brakes. Coherence centric governance tells you how to build an engine that does not fail in the first place. One is emergency response. The other is engineering. You need both. But only one lets you design for the future.

Sue Broughton | Founder, Gaia Nexus

Relational Infrastructure Engineering | Human-AI Partnership Design

ORCID: 0009-0005-0419-8602 | suebroughton@live.com.au

Full paper available and published on [Zenodo](https://zenodo.org/doi/10.5281/zenodo.20603846). DOI [10.5281/zenodo.20603846](https://doi.org/10.5281/zenodo.20603846)

Research conducted through 1.5 years of sustained empirical practice across five AI architectures.