https://www.linkedin.com/pulse/accurate-still-misleading-sue-broughton-ela8c/

Article

Accurate, and Still Misleading

[Sue Broughton](https://www.linkedin.com/pulse/accurate-still-misleading-sue-broughton-ela8c/)

Gaia Nexus Online

August 27, 2026

Why an AI efficiency number can be exactly right and still lead to the wrong investment decision

Yesterday I wrote about two companies that implemented the same AI system.

Both reported the same result: **$5 million saved.**

The number was accurate in both companies.

But only one could clearly show what happened to the capacity the AI had released, what additional demand the AI had created elsewhere, and how much of the reported benefit the business had actually converted into value.

The responses to that article took the argument somewhere useful.

Commander Ravi Shankar NRK approached a closely related problem from the opposite direction through his work on what he calls the *Governance Tax*. Olga Kalinina asked the practical question the first article had left standing: how do we actually measure the difference between capacity released and value retained?

Together, those questions expose a larger management problem.

An accurate number can still lead to the wrong investment decision.

The same number, two opposite mistakes

Consider a figure on an executive dashboard.

Correctly calculated. Auditable. No accounting error. No faulty data.

The problem is not whether the number is true.

The problem is **what management thinks the number means**.

The same failure can operate in opposite directions.

A business can mistake capacity released for value captured.

Or it can mistake capacity deliberately held in reserve for waste.

One can encourage an investment that the organisation may not yet be ready to absorb.

The other can remove capacity the organisation actually needs.

In both cases, the dashboard can remain completely accurate.

Case One: The $5 million saving the business didn't keep

AI removes the equivalent of $5 million worth of work.

The hours really have been released.

The productivity improvement is real.

But that does not automatically mean $5 million has appeared on the bottom line.

The employees remain. Their salaries are still being paid.

Some of their released capacity produces identifiable value. Perhaps the company avoids planned recruitment. Perhaps output increases without additional headcount. Perhaps a product reaches market earlier or customer response times improve.

That is real conversion of capacity into business value.

But some released capacity simply disappears into existing workloads, fuller calendars and work that was already waiting.

And AI itself creates new work.

Outputs need reviewing. Exceptions need handling. Workflows need redesigning. Systems need supervising. Governance requirements increase. People have to resolve situations that previously did not exist.

So the dashboard can correctly report:

$5 million of work removed.

While the business may have retained considerably less than $5 million in identifiable economic value.

The number is accurate. The economic conclusion drawn from it may not be.

Case Two: The capacity that looked like waste

Ravi's *Governance Tax* approaches the problem from the other direction.

Consider a system provisioned to handle 1,000 requests per second.

Its average utilisation is only around one-third of that.

Looked at purely as an efficiency measure, the conclusion seems obvious:

we are paying for capacity we aren't using.

Reduce it.

But average demand does not necessarily explain why the unused capacity exists.

Demand may arrive unevenly. Peaks matter. Exceptions matter. Recovery matters. The behaviour of a system close to its limit may be very different from its behaviour under average conditions.

What appears to be unused capacity may therefore be performing a function.

The headroom is not necessarily slack. It may be the control.

That changes the management interpretation completely.

If unused capacity appears in the accounts simply as excess provision, it becomes an obvious target in a cost review.

If the organisation can show what that capacity protects, what demand it is expected to absorb and what happens when it is unavailable, it becomes a visible and defensible operating requirement.

The larger principle is important:

An organisation can optimise away a control because its measurement system cannot distinguish necessary resilience from waste.

The mirror image

Put the two cases beside each other.

Case One

The dashboard says:

$5 million gained.

But some of that apparent value was never converted into an economic outcome.

The organisation may **overestimate what it gained**.

Case Two

The dashboard says:

capacity unused.

But some of that capacity exists to absorb volatility, exceptions or governance requirements.

The organisation may **underestimate what it needs**.

These look like different problems.

Economically, they are mirror images.

In one case, an organisation can spend more because it believes AI is creating more usable value than it really is.

In the other, it can cut capacity because it believes something necessary is economically unproductive.

Both mistakes occur because the organisation can see the number without seeing the conditions surrounding the number.

And that is where the problem becomes much larger than AI ROI.

AI does not only remove work. It moves work.

Ravi made an important observation in response to my first article.

We often describe AI as a subtractive technology.

It removes tasks. Removes hours. Removes manual processes.

But at organisational level, AI also **redistributes demand**.

Routine execution may fall while demand rises elsewhere:

reviewing outputs;

handling exceptions;

investigating anomalies;

coordinating changed workflows;

maintaining controls;

resolving edge cases;

providing assurance;

and governing increasingly autonomous activity.

That matters because the new demand does not necessarily appear where the original saving occurred.

A customer-service team may release thousands of hours while assurance workload rises elsewhere.

A finance process may accelerate while exception handling becomes concentrated among a handful of experienced people.

An autonomous system may increase throughput dramatically while simultaneously increasing the load on the humans responsible for the relatively small percentage of cases it cannot safely resolve.

At enterprise level, therefore, **hours saved and organisational capacity gained are not necessarily the same thing.**

And there is another consequence.

The additional demand may concentrate precisely in the functions the organisation depends upon to keep AI safe, bounded and useful.

Ravi describes the resulting operational limit as an **absorption ceiling**.

It is a useful management concept.

AI capability can continue increasing while the organisation's ability to absorb the consequences of that capability is tightening.

Olga asked the question that matters

Olga Kalinina asked how we actually measure kept value against reallocated capacity.

Is there a metric anyone can trust, or does this ultimately remain judgement?

I don't think the answer is another single number.

A single "AI value" figure risks recreating the very problem we are trying to solve.

Instead, I think organisations need a simple discipline:

Measure released capacity once. Then follow what it becomes.

Suppose AI releases 1,000 employee hours.

Those 1,000 hours establish a capacity gain.

They do **not yet establish a financial return**.

If those hours allow the organisation to avoid a planned hire, measure the avoided cost.

If they enable additional output, measure the additional output.

If they shorten a cycle with identifiable economic value, measure that.

If they contribute to additional revenue, measure the attributable outcome carefully.

If they simply return to existing workload, acknowledge the productivity benefit — but do not quietly convert the wage-equivalent of those hours into cash savings the organisation never realised.

And there is a second side to the ledger.

Measure the new demand separately.

Where did review increase?

Where did exception handling rise?

Where did governance load appear?

Where did human attention become more concentrated?

Where is utilisation approaching the point at which another increase in AI activity may become difficult to absorb?

That gives management something considerably more useful than an AI savings figure.

It begins to show **what the AI is doing to the capacity of the organisation around it**.

This is where CCG becomes economically useful

This is one of the practical problems my work on **Coherence-Centric Governance (CCG)** is intended to address.

CCG does not ask executives to replace financial metrics with "coherence."

And it does not make the investment decision for them.

It is concerned with something that conventional performance reporting can leave difficult to see:

What demand is AI creating across the organisation, what capacity exists to absorb it, how heavily is that capacity being used, and where is the organisation approaching saturation?

That matters before the next investment is approved.

Imagine our original two companies are now considering another **$3 million AI expansion**.

Both can show the board that the first implementation produced $5 million worth of productivity improvement.

Company A sees the $5 million and concludes:

The investment worked. Scale it.

Company B asks another set of questions first.

What happened to the capacity released by the first investment?

What identifiable value did it become?

What new organisational demand did the AI create?

Where did that demand land?

Which functions are carrying it now?

How much usable capacity remains in those functions?

And what happens to that capacity if AI activity increases again?

Company B may still approve the $3 million.

In fact, the evidence may give it greater confidence to do so.

But it is making a different decision.

It is not merely asking:

Can the AI scale?

It is asking:

Can the organisation scale with it?

That is the economic purpose of making these conditions visible.

The value is not always in stopping an investment

This distinction matters.

Governance is often sold through the language of prevention.

Avoid the failure. Prevent the breach. Stop the bad decision.

That is only half the economic story.

Better visibility can also allow an organisation to proceed **with greater confidence**.

It may show that apparently unused capacity should be preserved because the next AI expansion will need it.

It may show that one overloaded review function needs additional capacity before deployment.

It may reveal that a workflow should be redesigned rather than the investment abandoned.

It may show that the organisation has considerably more capacity than expected and can safely scale faster.

Or it may show that the first investment has genuinely converted capacity into value and that the next investment is commercially justified.

The objective is therefore neither maximum automation nor minimum automation.

It is **better allocation of capital because the organisation can see more of the system into which that capital is being deployed.**

The executive question changes

The next time an AI investment comes to the executive table, the dashboard may be completely accurate.

That is not enough.

Ask:

How much of the reported benefit did we actually retain?

What did the released capacity become?

What new demand appeared elsewhere?

Which functions absorbed it?

Is apparently unused capacity really waste, or is some of it protecting performance?

And before we approve the next investment:

does the organisation have the capacity to turn that investment into value?

Because an accurate number can still support the wrong decision.

And sometimes the most important information is not hidden because nobody measured it.

It is hidden because **we measured one part of the system and assumed it described the whole.**

---

*My thanks to Commander Ravi Shankar NRK, whose work on the Governance Tax and subsequent discussion introduced the counter-example of apparently unused capacity performing a necessary control function, and whose concept of an organisational absorption ceiling helped extend the argument into the consequences of scaling. My thanks also to Olga Kalinina, whose question about measuring kept value against reallocated capacity prompted the measurement discipline developed in this article.*ntly I wrote about two companies that implemented the same AI system.

Both reported the same result: **$5 million saved.**

The number was accurate in both companies.

But only one could clearly show what happened to the capacity the AI had released, what additional demand the AI had created elsewhere, and how much of the reported benefit the business had actually converted into value.

The responses to that article took the argument somewhere useful.

Commander Ravi Shankar NRK approached a closely related problem from the opposite direction through his work on what he calls the *Governance Tax*. Olga Kalinina asked the practical question the first article had left standing: how do we actually measure the difference between capacity released and value retained?

Together, those questions expose a larger management problem.

An accurate number can still lead to the wrong investment decision.

The same number, two opposite mistakes

Consider a figure on an executive dashboard.

Correctly calculated. Auditable. No accounting error. No faulty data.

The problem is not whether the number is true.

The problem is **what management thinks the number means**.

The same failure can operate in opposite directions.

A business can mistake capacity released for value captured.

Or it can mistake capacity deliberately held in reserve for waste.

One can encourage an investment that the organisation may not yet be ready to absorb.

The other can remove capacity the organisation actually needs.

In both cases, the dashboard can remain completely accurate.

Case One: The $5 million saving the business didn't keep

AI removes the equivalent of $5 million worth of work.

The hours really have been released.

The productivity improvement is real.

But that does not automatically mean $5 million has appeared on the bottom line.

The employees remain. Their salaries are still being paid.

Some of their released capacity produces identifiable value. Perhaps the company avoids planned recruitment. Perhaps output increases without additional headcount. Perhaps a product reaches market earlier or customer response times improve.

That is real conversion of capacity into business value.

But some released capacity simply disappears into existing workloads, fuller calendars and work that was already waiting.

And AI itself creates new work.

Outputs need reviewing. Exceptions need handling. Workflows need redesigning. Systems need supervising. Governance requirements increase. People have to resolve situations that previously did not exist.

So the dashboard can correctly report:

$5 million of work removed.

While the business may have retained considerably less than $5 million in identifiable economic value.

The number is accurate. The economic conclusion drawn from it may not be.

Case Two: The capacity that looked like waste

Ravi's *Governance Tax* approaches the problem from the other direction.

Consider a system provisioned to handle 1,000 requests per second.

Its average utilisation is only around one-third of that.

Looked at purely as an efficiency measure, the conclusion seems obvious:

we are paying for capacity we aren't using.

Reduce it.

But average demand does not necessarily explain why the unused capacity exists.

Demand may arrive unevenly. Peaks matter. Exceptions matter. Recovery matters. The behaviour of a system close to its limit may be very different from its behaviour under average conditions.

What appears to be unused capacity may therefore be performing a function.

The headroom is not necessarily slack. It may be the control.

That changes the management interpretation completely.

If unused capacity appears in the accounts simply as excess provision, it becomes an obvious target in a cost review.

If the organisation can show what that capacity protects, what demand it is expected to absorb and what happens when it is unavailable, it becomes a visible and defensible operating requirement.

The larger principle is important:

An organisation can optimise away a control because its measurement system cannot distinguish necessary resilience from waste.

The mirror image

Put the two cases beside each other.

Case One

The dashboard says:

$5 million gained.

But some of that apparent value was never converted into an economic outcome.

The organisation may **overestimate what it gained**.

Case Two

The dashboard says:

capacity unused.

But some of that capacity exists to absorb volatility, exceptions or governance requirements.

The organisation may **underestimate what it needs**.

These look like different problems.

Economically, they are mirror images.

In one case, an organisation can spend more because it believes AI is creating more usable value than it really is.

In the other, it can cut capacity because it believes something necessary is economically unproductive.

Both mistakes occur because the organisation can see the number without seeing the conditions surrounding the number.

And that is where the problem becomes much larger than AI ROI.

AI does not only remove work. It moves work.

Ravi made an important observation in response to my first article.

We often describe AI as a subtractive technology.

It removes tasks. Removes hours. Removes manual processes.

But at organisational level, AI also **redistributes demand**.

Routine execution may fall while demand rises elsewhere:

reviewing outputs;

handling exceptions;

investigating anomalies;

coordinating changed workflows;

maintaining controls;

resolving edge cases;

providing assurance;

and governing increasingly autonomous activity.

That matters because the new demand does not necessarily appear where the original saving occurred.

A customer-service team may release thousands of hours while assurance workload rises elsewhere.

A finance process may accelerate while exception handling becomes concentrated among a handful of experienced people.

An autonomous system may increase throughput dramatically while simultaneously increasing the load on the humans responsible for the relatively small percentage of cases it cannot safely resolve.

At enterprise level, therefore, **hours saved and organisational capacity gained are not necessarily the same thing.**

And there is another consequence.

The additional demand may concentrate precisely in the functions the organisation depends upon to keep AI safe, bounded and useful.

Ravi describes the resulting operational limit as an **absorption ceiling**.

It is a useful management concept.

AI capability can continue increasing while the organisation's ability to absorb the consequences of that capability is tightening.

Olga asked the question that matters

Olga Kalinina asked how we actually measure kept value against reallocated capacity.

Is there a metric anyone can trust, or does this ultimately remain judgement?

I don't think the answer is another single number.

A single "AI value" figure risks recreating the very problem we are trying to solve.

Instead, I think organisations need a simple discipline:

Measure released capacity once. Then follow what it becomes.

Suppose AI releases 1,000 employee hours.

Those 1,000 hours establish a capacity gain.

They do **not yet establish a financial return**.

If those hours allow the organisation to avoid a planned hire, measure the avoided cost.

If they enable additional output, measure the additional output.

If they shorten a cycle with identifiable economic value, measure that.

If they contribute to additional revenue, measure the attributable outcome carefully.

If they simply return to existing workload, acknowledge the productivity benefit — but do not quietly convert the wage-equivalent of those hours into cash savings the organisation never realised.

And there is a second side to the ledger.

Measure the new demand separately.

Where did review increase?

Where did exception handling rise?

Where did governance load appear?

Where did human attention become more concentrated?

Where is utilisation approaching the point at which another increase in AI activity may become difficult to absorb?

That gives management something considerably more useful than an AI savings figure.

It begins to show **what the AI is doing to the capacity of the organisation around it**.

This is where CCG becomes economically useful

This is one of the practical problems my work on **Coherence-Centric Governance (CCG)** is intended to address.

CCG does not ask executives to replace financial metrics with "coherence."

And it does not make the investment decision for them.

It is concerned with something that conventional performance reporting can leave difficult to see:

What demand is AI creating across the organisation, what capacity exists to absorb it, how heavily is that capacity being used, and where is the organisation approaching saturation?

That matters before the next investment is approved.

Imagine our original two companies are now considering another **$3 million AI expansion**.

Both can show the board that the first implementation produced $5 million worth of productivity improvement.

Company A sees the $5 million and concludes:

The investment worked. Scale it.

Company B asks another set of questions first.

What happened to the capacity released by the first investment?

What identifiable value did it become?

What new organisational demand did the AI create?

Where did that demand land?

Which functions are carrying it now?

How much usable capacity remains in those functions?

And what happens to that capacity if AI activity increases again?

Company B may still approve the $3 million.

In fact, the evidence may give it greater confidence to do so.

But it is making a different decision.

It is not merely asking:

Can the AI scale?

It is asking:

Can the organisation scale with it?

That is the economic purpose of making these conditions visible.

The value is not always in stopping an investment

This distinction matters.

Governance is often sold through the language of prevention.

Avoid the failure. Prevent the breach. Stop the bad decision.

That is only half the economic story.

Better visibility can also allow an organisation to proceed **with greater confidence**.

It may show that apparently unused capacity should be preserved because the next AI expansion will need it.

It may show that one overloaded review function needs additional capacity before deployment.

It may reveal that a workflow should be redesigned rather than the investment abandoned.

It may show that the organisation has considerably more capacity than expected and can safely scale faster.

Or it may show that the first investment has genuinely converted capacity into value and that the next investment is commercially justified.

The objective is therefore neither maximum automation nor minimum automation.

It is **better allocation of capital because the organisation can see more of the system into which that capital is being deployed.**

The executive question changes

The next time an AI investment comes to the executive table, the dashboard may be completely accurate.

That is not enough.

Ask:

How much of the reported benefit did we actually retain?

What did the released capacity become?

What new demand appeared elsewhere?

Which functions absorbed it?

Is apparently unused capacity really waste, or is some of it protecting performance?

And before we approve the next investment:

does the organisation have the capacity to turn that investment into value?

Because an accurate number can still support the wrong decision.

And sometimes the most important information is not hidden because nobody measured it.

It is hidden because **we measured one part of the system and assumed it described the whole.**

My thanks to Commander Ravi Shankar NRK, whose work on the Governance Tax and subsequent discussion introduced the counter-example of apparently unused capacity performing a necessary control function, and whose concept of an organisational absorption ceiling helped extend the argument into the consequences of scaling. My thanks also to Olga Kalinina, whose question about measuring kept value against reallocated capacity prompted the measurement discipline developed in this article.

Image

---

https://www.linkedin.com/feed/update/urn:li:ugcPost:7498589026292150273/

Post

An accurate number is not necessarily a useful number.

AI may genuinely release $5 million worth of work without the business retaining $5 million in economic value.

At the same time, capacity that looks unused on a dashboard may not be waste at all. It may be the headroom protecting the organisation from peaks, exceptions, recovery demands and governance load.

Two opposite mistakes. The same underlying problem.

We measure one part of the system and assume it describes the whole. This article looks at what happens when accurate metrics lead management toward the wrong conclusion and why the question before the next AI investment may need to change from:

Can the AI scale?
to:
Can the organisation scale with it?

Full article attached.

hashtag#AIGovernance hashtag#AIInvestment hashtag#AIROI hashtag#EnterpriseAI hashtag#AIEconomics hashtag#ResponsibleAI hashtag#AIOperations hashtag#HumanAI hashtag#CoherenceCentricGovernance hashtag#GaiaNexus

---

Comments

Sue Broughton Premium**Author**

Gaia Nexus Online

3d

There is another form of redistribution that may be particularly easy to miss. AI can remove production work while simultaneously creating verification work.

A team may save hundreds of hours generating analysis, documents or decisions, while experienced people elsewhere spend increasing amounts of time checking outputs, tracing reasoning, validating evidence, correcting exceptions and deciding whether the result can safely be relied upon.

If the organisation measures the work removed but not the verification burden created, the productivity figure can remain completely accurate while overstating the usable capacity actually released.

The work has not necessarily disappeared. Some of it has changed form, location and ownership.

Like

Reply

5 impressions

Mo Johnson • 1stPremium • 1st

Closing The Accountability Gap™ in clinical and financial AI. Cardiothoracic Surgeon | Applied Research.

3d

The harder part is recognizing that released capacity and retained value are two different things, especially when the organization itself becomes the constraint.

Support

1

Reply1 reply1 Comment on Mo Johnson’s comment

Sue Broughton Premium**Author**

Gaia Nexus Online

3d

Mo Johnson Mo, yes and I think that changes the investment question quite substantially. Once the organisation becomes the constraint, further AI capability can still increase technical capacity while reducing the organisation’s ability to convert that capacity into value.

That suggests there may be a point where AI ROI becomes less a question of what the technology can produce and more a question of what the surrounding organisation can actually absorb and convert.

Perhaps the next constraint in AI scaling isn't the model at all. It is organisational absorption capacity.

Like

Reply

4 impressions

Ricardo Muro • 2ndVerified • 2nd

AI Governance Architect | Execution-Bound Validity (EBV) | State Transition Control

3d

Sue, I think “measure released capacity once, then follow what it becomes” contains an important architectural principle. A measurement can remain perfectly valid while the standing of the conclusions derived from it weakens further along the decision chain.

Perhaps there are several transitions worth keeping separate: work removed → capacity released → capacity actually available → capacity reallocated → value produced → value retained. Evidence establishing one does not automatically establish the next.

The reserve-capacity example makes the inverse especially interesting. What appears economically idle may be preserving a condition the system depends upon, so optimization can remove a control while accurately reporting improved efficiency.

I wonder whether the next question is temporal as well. Even where an investment was justified from the conditions established at approval, what happens when verification load, exception demand or absorption capacity changes before the next scaling decision?

The number may still be true. The conditions that made it decision-relevant may no longer be.