

<https://www.linkedin.com/pulse/when-ai-changes-organisation-around-sue-broughton-tr5mc/?trackingId=Dw2D0qrK5CQff%2Fy%2BSxjVoQ%3D%3D>

Article

When AI Changes the Organisation Around It



Sue Broughton

Gaia Nexus Online

August 21, 2026

Gaia Nexus Online

August 21, 2026

The next AI governance question may not be what the AI is doing, but what sustained AI use is doing to the organisation.

Imagine an organisation where AI adoption appears to be going extremely well.

Decision times are falling. Productivity is increasing. Human overrides are becoming less frequent. Escalations are declining. Review processes are getting faster. Employees are increasingly comfortable relying on AI-supported decisions.

By many conventional measures, this looks like successful AI adoption.

But there is another possibility worth considering.

What if some of those same indicators could also accompany a gradual reduction in independent human judgement, challenge behaviour or willingness to intervene?

The difficulty is that **efficiency and erosion can sometimes produce remarkably similar signals.**

When Improvement and Erosion Look Alike

Consider something as simple as declining escalation rates.

There are several reasons this might happen.

The AI may have become more accurate. Employees may have become more skilled at using it. Processes may have improved and fewer cases may genuinely require escalation.

All of those would be positive developments.

But escalation might also decline because people have become increasingly accustomed to accepting AI recommendations.

The observable result is the same.

Escalations decrease.

The meaning is entirely different.

The same ambiguity can appear elsewhere.

Fewer human overrides might indicate greater AI accuracy, or increasing human deference.

Faster review times might indicate improved workflow efficiency, or increasingly superficial review.

Less double checking might reflect greater system reliability, or declining independent verification.

Greater trust might represent healthy Human AI collaboration, or growing dependency.

None of these behaviours should automatically be interpreted as positive or negative.

The governance challenge is determining **what they mean in context and how that meaning changes over time.**

AI Can Change More Than the Work

Much of the discussion around AI adoption understandably focuses on the capabilities and behaviour of AI systems.

Is the system accurate?

Is it safe?

Is it compliant?

Can its outputs be explained?

Can its decisions be audited?

These remain important questions.

But as AI becomes embedded in everyday organisational activity, another question emerges:

What is sustained interaction with AI doing to the organisation around it?

People adapt to the systems they use.

They learn when to trust them, when to check them, when to challenge them, when to escalate and when to defer.

Initially these may simply be individual behaviours.

Repeated across thousands of interactions, however, behaviours can become routines.

Routines can become norms.

Norms can eventually reshape organisational pathways.

The organisation may gradually change how it thinks, decides, checks, challenges, escalates and allocates trust.

And none of this requires an AI system to fail.

When Successful AI Creates a Different Governance Question

There is an interesting paradox here.

The more consistently useful an AI system becomes, the more rational it may appear for humans to rely upon it.

As confidence grows, checking may decrease.

As recommendations prove reliable, overrides may become less frequent.

As employees become accustomed to AI assistance, some forms of independent judgement may be exercised less often.

That does not necessarily mean capability has disappeared.

But it raises an important distinction:

Capability being formally present is not the same as capability being regularly exercised.

An organisation may still employ highly qualified analysts, engineers, accountants, clinicians, managers or other professionals.

Their credentials have not changed.

Their authority may not have changed.

But if AI increasingly performs parts of the reasoning, interpretation or assessment that those people previously exercised themselves, how does the organisation know whether the underlying human capability remains readily available when circumstances genuinely require it?

This becomes particularly important where human intervention is itself part of the organisation's governance architecture.

A control that depends upon human judgement ultimately depends upon that judgement remaining available.

The Effects May Be Longitudinal

Traditional governance is often designed to detect events.

An error occurs.

A threshold is breached.

A control fails.

An incident is reported.

Someone investigates.

But organisational adaptation to AI may not look like an event.

There may be no identifiable moment when something went wrong.

Instead, there may be hundreds of thousands of individually reasonable decisions:

This output looks right.

We don't need the second check this time.

The system is usually accurate.

There is probably no reason to escalate this.

I'll accept the recommendation.

Each decision may be perfectly defensible.

The interesting question is what happens when these behaviours accumulate over months or years.

The relevant unit of observation may therefore need to extend beyond individual AI interactions toward **patterns of organisational adaptation over time.**

I have been thinking about this as **Enterprise Pathway Drift**: gradual changes in organisational decision pathways, checking behaviours, escalation patterns, challenge practices, trust allocation and human capability that emerge through sustained interaction with AI.

Drift does not necessarily mean deterioration.

Some pathways should change.

AI should remove unnecessary work, improve decisions and make organisations more effective.

The challenge is distinguishing **healthy adaptation from gradual erosion**.

The Organisation May Already Hold Some of the Evidence

Interestingly, organisations may not always need entirely new forms of data to begin asking these questions.

Evidence may already exist across operational systems.

Escalation histories may show how frequently people seek additional judgement.

Override records may reveal changing patterns of human disagreement with automated recommendations.

Review logs may show whether verification behaviour is changing.

Exception records may reveal emerging workarounds.

Workflow histories may show where decision pathways are compressing or changing.

Survey and qualitative evidence may reveal changing patterns of trust, confidence and willingness to challenge.

Individually, none of these measures tells us whether an organisation is becoming more capable or less capable.

Together, and observed longitudinally, they may tell a much richer story.

Perhaps the missing element is not always more data.

Perhaps sometimes it is **a different question asked of the data we already have**.

The Visibility Problem

There is another difficulty.

The people who notice organisational change first may not be the people with authority to respond to it.

A junior analyst may notice that colleagues rarely question AI recommendations anymore.

A frontline employee may see an unofficial workaround becoming normal practice.

A supervisor may notice that a review step still formally exists but is increasingly treated as a formality.

These observations can be extremely valuable.

But organisations often separate **visibility from authority**.

Senior leaders have authority but frequently see aggregated information.

Frontline workers have direct visibility but limited authority.

This creates an important governance challenge:

How does an organisation allow observations from the edges of the system to become legitimate evidence without confusing observation with decision authority?

The answer cannot simply be to give everyone authority over everything.

But neither should valuable observations disappear because the person who noticed the change lacks organisational power.

Governance may increasingly require better pathways between **observation, evidence, interpretation and authorised human decision-making**.

A Different Kind of AI Governance Question

None of this replaces existing AI governance.

Accuracy matters.

Safety matters.

Compliance matters.

Bias, security, explainability and technical reliability matter.

But increasingly capable AI may require us to observe another system as well:

the human organisation adapting around it.

That means asking questions that conventional AI performance measures may not answer.

Are people challenging AI differently than they did a year ago?

Are escalation pathways still being used?

Is independent verification increasing or decreasing?

Are different teams developing different relationships with the same AI systems?

Are experienced and junior employees adapting in the same way?

Are human capabilities being strengthened, displaced or simply exercised less frequently?

And perhaps most importantly:

Would the organisation recognise the difference between becoming more efficient and becoming more dependent?

We may discover that some of the most consequential effects of AI adoption do not initially appear as failures.

Some may appear as improvements.

That is why longitudinal observation matters.


As AI becomes increasingly embedded in organisational decision-making, governance may need to look beyond whether AI systems continue to perform well.

We may also need to understand what prolonged interaction with those systems is gradually doing to the humans, behaviours and decision pathways around them.

The question I think deserves more attention is therefore relatively simple:

If AI is changing how an organisation thinks, decides, checks, challenges, escalates and trusts, how would we know whether those changes represent healthy adaptation or gradual erosion?

Image



Post

<https://www.linkedin.com/feed/update/urn:li:ugcPost:7496483743843397633/>



Sue Broughton • YouPremium • YouGaia Nexus
 OnlineGaia Nexus Online 1w • 1 week ago • Visible to anyone on or off LinkedIn

What if some of the signals we interpret as successful AI adoption could also be signals of something else?

Fewer overrides.
 Fewer escalations.
 Faster reviews.

Greater trust.
Less double checking.

These may indicate that AI is working extremely well.

But over time, they could also accompany declining independent judgement, reduced challenge behaviour or increasing human dependency.

And from the outside, the two can look remarkably similar.

This raises a governance question I think deserves much more attention: As AI becomes embedded across an organisation, how do we distinguish healthy adaptation from gradual erosion?

My latest article explores what I call Enterprise Pathway Drift and why the organisation adapting around AI may increasingly need to become part of what AI governance observes.

Full article attached

[hashtag#AIGovernance](#) [hashtag#EnterpriseAI](#)
[hashtag#ResponsibleAI](#) [hashtag#HumanAI](#) [hashtag#AIagents](#)
[hashtag#OrganisationalGovernance](#) [hashtag#HumanReadiness](#)
[hashtag#EnterprisePathwayDrift](#)

Comments



[Jay VonSpreck](#) • 1stVerified • 1st

[Founder | Quantum Competitive Intelligence™ Architect | Risk Command Architecture™ Executive Intelligence Systems \(EIS\) Execution Governance | Balance Sheet Outcomes](#)

1w

Sue, this is an important distinction. A reduction in overrides or escalations can look like improved performance while the execution system producing those numbers is quietly changing.

The governance question therefore cannot stop at whether the AI is behaving as intended. It has to include whether the conditions surrounding execution remain governable as humans adapt to the system.

Fewer interventions may represent healthy adaptation. They may also represent migration of authority, judgment, challenge behavior, or intervention capacity.

The number alone cannot tell us which occurred. That is why we treat the interacting conditions around execution—not merely the technology—as part of the governing object. The trajectory matters.

Support

2

Reply 2 replies 2 Replies on Jay VonSpreck's comment

See previous replies



Jay VonSpreck • 1stVerified • 1st

[Founder | Quantum Competitive Intelligence™ Architect | Risk Command Architecture™ Executive Intelligence Systems \(EIS\) Execution Governance | Balance Sheet Outcomes](#)

1w

Sue Broughton Sue, exactly. And I think this is where trajectory and intervention rights become inseparable.

Observation can reveal that authority, judgment or challenge behavior is migrating. But observation alone does not make that migration governable.

The architecture must also establish who has the authority to intervene, under what conditions, and while the trajectory can still be changed.

That creates an important distinction: migration itself is not necessarily

failure. The governing question is whether the movement is intentional, observable, authorized and reversible—or whether consequential change is occurring outside the enterprise’s ability to govern it.

That may be one of the most important human-AI governance questions emerging right now.

Like

2

Reply



Farah Ahmed • 1stVerified • 1st

[Director, Sales & CX Transformation, StorageVault Canada | Leading in the Agentic Era - execution compresses, judgment does not | 272 locations | Writing “The Operating System”](#)

1w

We hit this exact question running AI in production. The tell we settled on: silence is only a good signal if people can still explain the overrides they're not making. So we periodically ask someone to walk through a case the system handled and defend or challenge the call. If they can, the trust is earned. If they can't, the quiet was atrophy wearing trust's clothes.

Support

2

Reply1 reply1 Comment on Farah Ahmed’s comment



Sue Broughton PremiumAuthor

[Gaia Nexus Online](#)

1w

Farah Ahmed Farah, atrophy wearing trust's clothes, captures the risk beautifully. And I particularly like that your test examines the overrides that aren't happening rather than treating silence itself as evidence of successful alignment.

I wonder whether the next layer is not only whether someone can explain the system's call, but whether they can still independently reconstruct the judgement, identify what would have changed it, and meaningfully disagree with it.

That might help distinguish retained understanding from retrospective agreement. What you are describing feels like a very practical way of testing whether trust has remained earned rather than quietly becoming dependence.

Like

1

Reply

22 impressions



Ricardo Muro • 2ndVerified • 2nd

[AI Governance Architect | Execution-Bound Validity \(EBV\) | State Transition Control](#)

1w

Sue, I think the distinction between a capability being formally present and remaining practically available is especially important. If human judgement or intervention is part of the governance architecture, then the continuing viability of that capability may itself need evidence.

Declining overrides or escalations cannot establish that either way. They may reflect better system performance, healthy adaptation, increasing

deference, changing incentives, or capability being exercised less often.

I wonder whether the harder longitudinal test is therefore not simply whether people are still in the loop, but whether they can still independently reconstruct the judgement when required: identify the relevant evidence, reach a conclusion without inheriting the system's answer, explain what would change that conclusion, and meaningfully challenge the recommendation.

That would make human oversight less a static role assignment and more a capability whose state can change over time.

Your Enterprise Pathway Drift framing raises an interesting governance question here: if the organisation itself is adapting around the AI, how do we establish that the controls we originally relied upon remain viable in the organisation that now exists?

Support

2

Reply1 reply1 Comment on Ricardo Muro's comment



[Sue Broughton](#) [Premium](#)**Author**

[Gaia Nexus Online](#)

1w

[Ricardo Muro](#) Ricardo, yes. I think your final question goes directly to the heart of Enterprise Pathway Drift. The control may remain formally unchanged while the organisation around it has changed sufficiently that the conditions which originally made that control viable no longer exist in the same form. I particularly like your idea of independently reconstructing the judgement. That begins to provide evidence of practical availability without reducing human capability to a static score. Can the person still identify the evidence, form an independent view, explain what would alter it and challenge the system where necessary?

It also raises another longitudinal question for me. Perhaps we should not only preserve evidence that a control was valid when introduced, but periodically ask whether the organisational and human conditions upon which its validity depended are still present.

In that sense, Enterprise Pathway Drift is not simply asking whether the AI has changed. It is asking whether the system of governance surrounding the AI has changed underneath controls that still appear intact.